

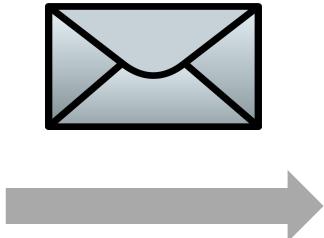
# Introduction to Kafka

# Outline

---

- Producers
  - What is a producer?
  - The producer API
- What is a Consumer?
  - What is a consumer and a consumer group?
  - The consumer API
- Anatomy of messages

# What is a Producer?

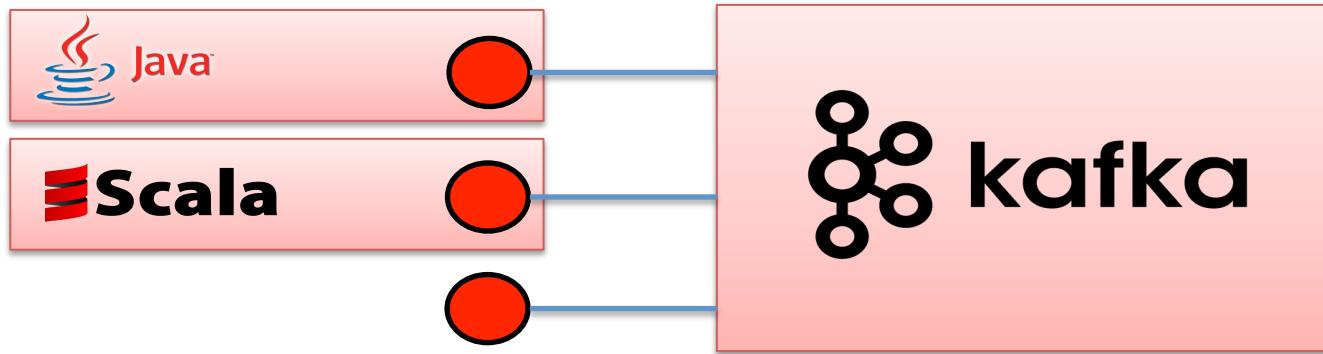


kafka

## ➤ Producer

- Producers produces the data sent to the Kafka clusters
  - Sent via topics
  - Directly involved in load-balancing
  - Controls the resiliency of messages

# Kafka APIs



- Kafka ships with built in client APIs for developers to use with applications
  - Kafka ships with a Java client that is recommended
  - Legacy Scala clients are still included
  - Kafka also includes a **binary wire protocol**
  - Many tools in other languages that implement this wire protocol

# The Java API

Generic sender where:  
K = Type of key  
V = Type of message

Constructor takes a configuration  
(mostly a hashmap of options)

Send a messages (with or without  
callbacks)

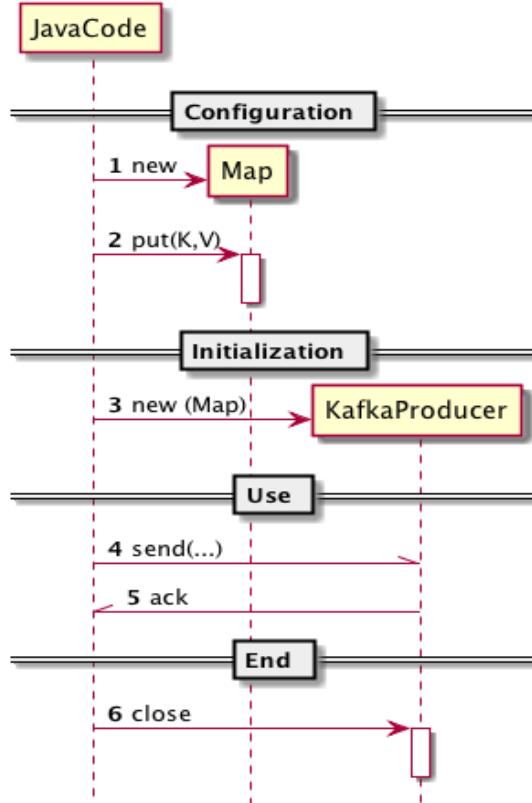
Get metrics for this producer

org.apache.kafka.clients.producer

KafkaProducer<K,V>

KafkaProducer(config: Properties) send(ProducerRecord<K,V>):  
Future<RecordMetaData>  
send(ProducerRecord<K,V>, Callback): Future<...>  
flush()  
metrics(): Map<MetricName, ? extends Metric>  
close()

# Java API Behavior



```
// Configuration  
Properties kp = new Properties();  
kp.put("bootstrap.servers",  
       "mybroker1:9092,mybroker2:9092");  
kp.put("key.serializer", "...");
```

```
// Initialization  
KafkaProducer<String, String> producer =  
    new KafkaProducer<String, String>(kp);
```

```
// Use  
Future<...> f = producer.send(...);  
... f.get(); // when acked
```

```
// End  
producer.close();
```

# Creating a Kafka Producer

---

- Constructing a Kafka producer requires 3 mandatory properties
  - `bootstrap.servers` – list of host:port pairs of Kafka brokers. This doesn't have to include all brokers in the cluster as the producer will query about additional brokers. It is recommended to include at least 2 in case one broker goes down
  - `key.serializer` – should be set to a class that implements the `Serializer` interface that will be used to serialize **keys**
  - `value.serializer` - should be set to a class that implements the `Serializer` interface that will be used to serialize **values**

# The ProducerRecord

Generic record where:  
K = message key  
V = Type of message

org.apache.kafka.clients.producer

ProducerRecord<K,V>

Message Key is optional

Partition Key is optional

key: K

value: V

topic: String

partition: Integer

ProducerRecord(topic: String, partition: Integer, key: K, value:V)

ProducerRecord(topic: String, key: K, value: V)

ProducerRecord(topic: String, value: V)

# Sending a Message

```
ProducerRecord<String, String> record =  
    new ProducerRecord<String, String>(  
        "someTopic", "someKey", "someValue");  
  
producer.send(record, new Callback() {  
    public void onCompletion(  
        RecordMetadata metadata, Exception e) {  
        if(e != null) e.printStackTrace();  
        System.out.println(  
            "Offset: " + metadata.offset());  
    }  
});
```

# Producer Controls Message Guarantees

---

- As a producer you can determine what guarantees you want Kafka to give you when sending a message
  - Controlled by acknowledgement
- Different use cases require different guarantees
  - Web page clicks log → Don't care if I lose a few messages
  - Credit card payment → I want best possible guarantee
- Kafka provides options
  - The better guarantee, the lower latency

# Producer API

---

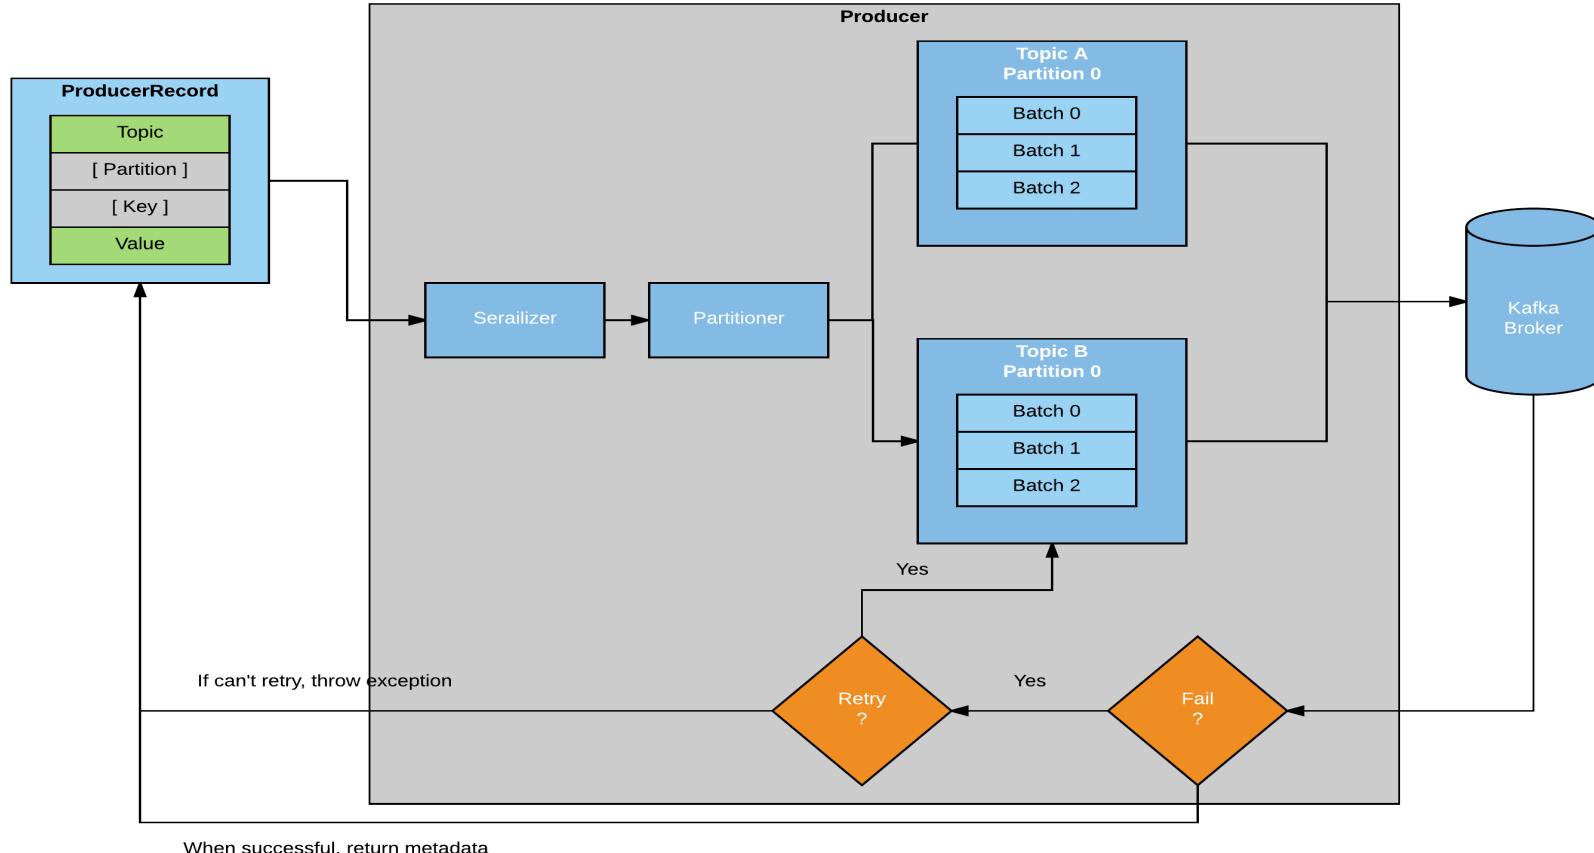
- Different use case requirements will influence the way the producer API is used to write messages to Kafka and its configuration
- Three primary ways of sending messages
  - Fire-and-forget – send a message and don't really care if it arrived successfully or not. Most of the time it will arrive successfully but it's possible that some messages will get lost
  - Synchronous send – message is sent and a Future object is returned which can be used to see if the send() was successful
  - Asynchronous send – the send() method has a callback function which is triggered when a response is received from the Kafka broker

# Acknowledgement of Messages

---

- No ack (0)
  - Kafka will most likely receive the message
  - Producer will not wait for any reply from the broker before assuming the message was sent successfully
  - If something goes wrong, producer will not know and message is lost
  - Because producer is not waiting for a response, high throughput can be achieved
- Ack from N replicas (1..N)
  - A message is not considered consumed by the Kafka cluster unless N replicas holding the message has acknowledged
- Ack from all replicas (-1)
  - Every replica must acknowledge the message

# Producer Overview



# Serialization

---

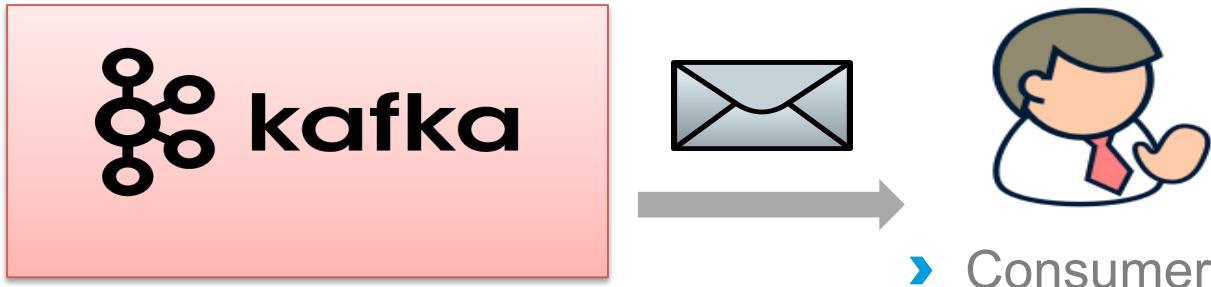
- Kafka messages are byte arrays (to Kafka)
  - Key → Array of bytes
  - Value → Array of bytes
- The Java API allows you to pass any object as key or value
  - Makes the code readable, but...
  - ... requires serializes and deserializes
- Kafka includes an interface for this  
`org.apache.kafka.common.serialization.Serializer`

# Built-in Serializers

---

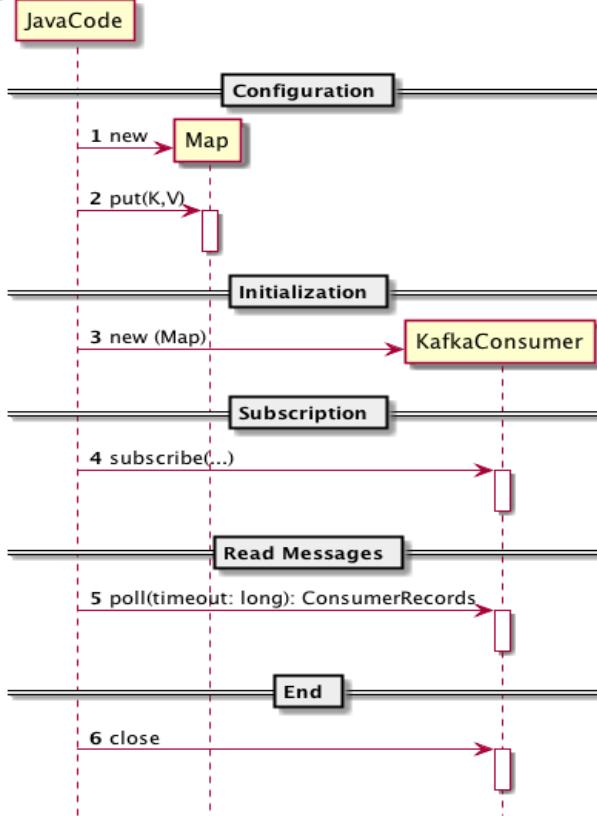
- Kafka includes serializers for common types:
  - ByteArraySerializer
  - StringSerializer
  - IntegerSerializer
  - ...
- Most organizations settle on some standard serialization strategy
  - JSON, XML, Apache Avro, Protobuf

# Consumers and Consumer Groups



- Applications that read data from Kafka are consumers
  - Subscribes to topics
  - Use KafkaConsumer to read messages from these topics
  - Kafka consumers are usually part of a *consumer group*
  - **The main way consumption of data from a Kafka topic is scaled is by adding more consumers to a consumer group**

# Java API Behavior



```
// Configuration
Properties kp = new Properties();
kp.put("bootstrap.servers",
       "mybroker1:9092,mybroker2:9092");
kp.put("key.deserializer", "...");

// Initialization
KafkaConsumer<...> consumer=
    new KafkaConsumer<...>(kp);

// Subscription
consumer.subscribe("interesting.*");

// Read messages
ConsumerRecords<...> records = consumer.poll(100);
for (ConsumerRecord<...> cr : records) {
    // cr.value(); cr.key(); cr.offset();
}

// End
consumer.close()
```

# The Java API

Constructor takes a configuration  
(mostly a hashmap of options)

Multiple ways to subscribe.  
Subscribe by list, wildcard, etc.

Read messages from Kafka

Multiple (sync/async) ways to  
confirm reception to consumer  
groups

A set of other methods such as: metrics(), pause(...), assign(...), close(), etc

Generic sender where:  
K = Type of key  
V = Type of message

org.apache.kafka.clients.consumer

KafkaConsumer<K,V>

KafkaConsumer(config: Properties)

subscribe(...)

poll(timeout: long): ConsumerRecords<K,V>

commit(...)

...

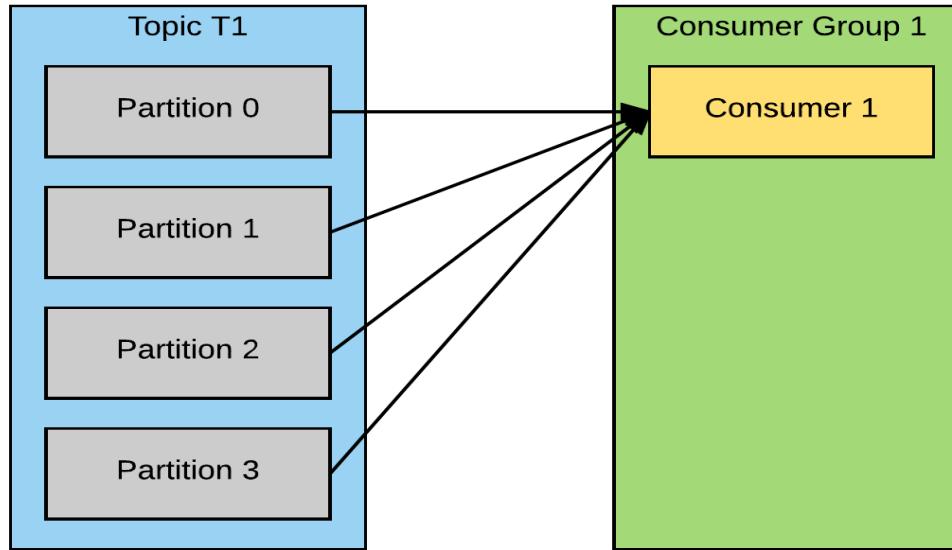
# How to Use the API?

---

- The ***org.apache.kafka.clients.consumer.KafkaConsumer*** acts as a proxy for the consumer
- Some key issues to resolve
  - Setup of consumer groups
  - Which topics to subscribe to
  - Which partitions to subscribe to (optional)
  - Manual or automatic offset management
  - Multi-threaded or single-threaded consumption

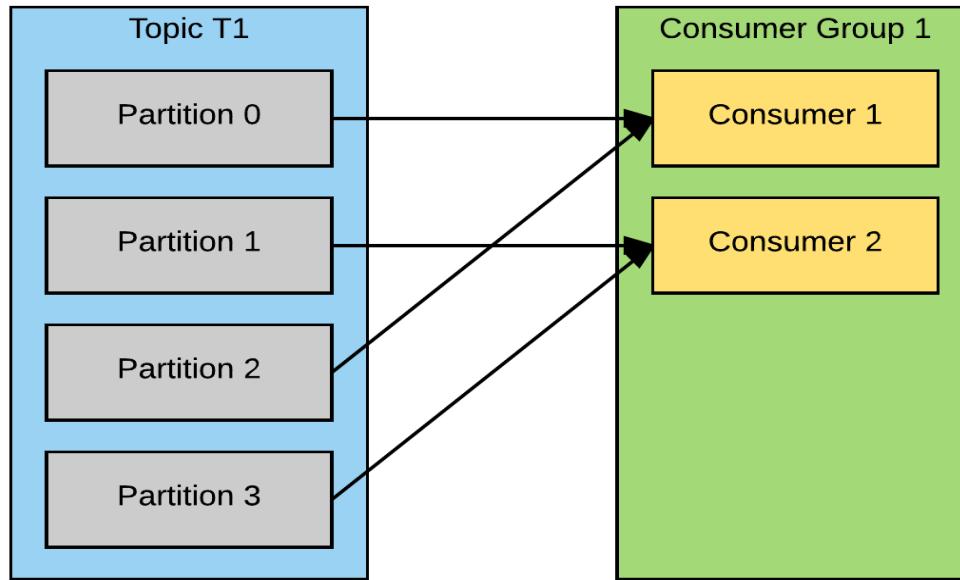
# Consumer Group

- The consumer will receive all messages from all four partitions



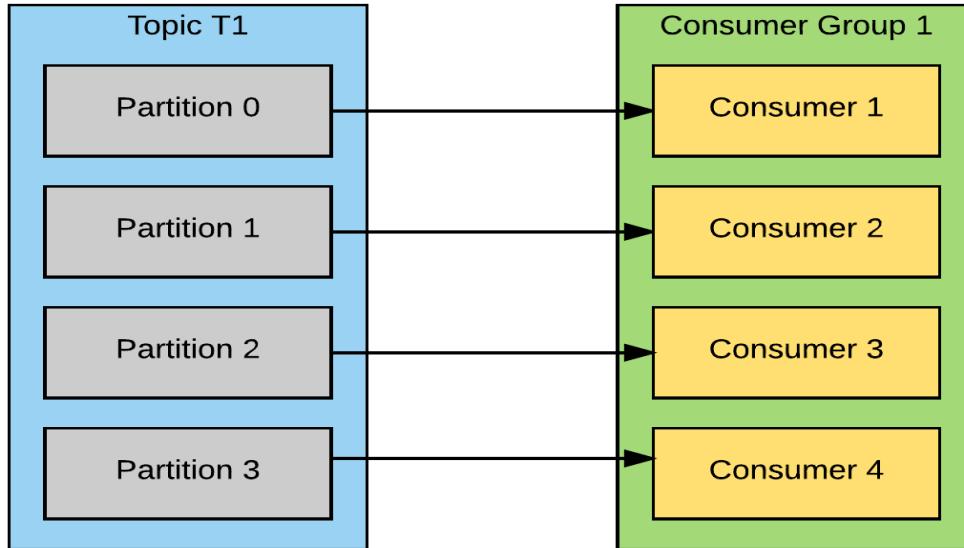
# Consumer Group

- Each consumer will only get messages from two partitions



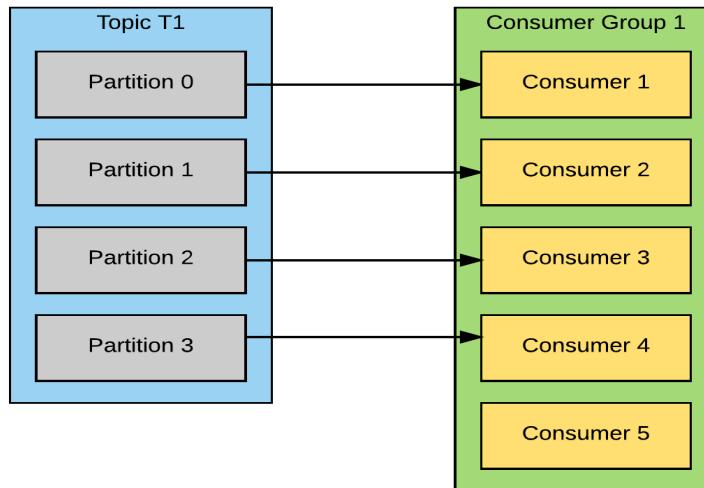
# Consumer Group

- If the consumer group has the same number of consumers as partitions, each will read messages from a single partition



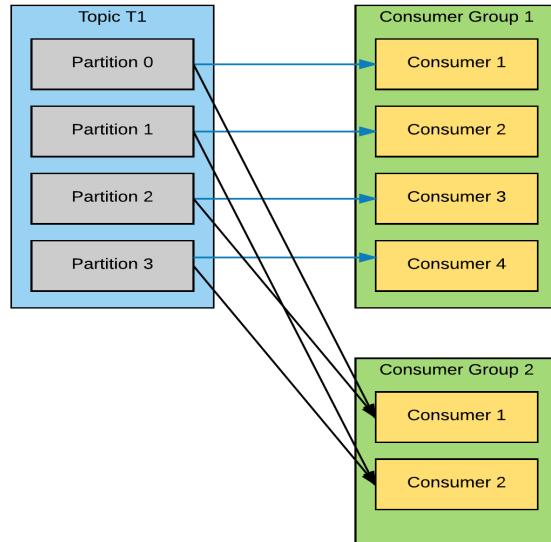
# Consumer Group

- If there are more consumers than partitions for a topic, some consumers will be idle and receive no messages
- Create topics with a large number of partitions to allow adding more consumers when load increases



# Multiple Consumer Groups

- One of the main design goals of Kafka was to allow multiple applications the ability to read data from the same topic
- Make sure each application has its own consumer group for this purpose



# Partition Rebalance

---

- Consumers in a consumer group share ownership of the partitions in the topics they subscribe to
- When a new consumer is **added** to the group, it consumes messages from partitions which were previously consumed by another consumer
- When a consumer **leaves** the group (shuts down, crashes, etc), the partitions it used to consume will be consumed by one of the remaining consumers
- Reassignment of partitions to consumers can also happen when topics are modified – an administrator adds new partitions, for example

# Partition Rebalance

---

- The event in which partition ownership is moved from one consumer to another is called a *rebalance*
- During a rebalance, consumers can't consume messages!
- Steps can be taken to safely handle rebalances and avoid unnecessary rebalances

# Creating a Kafka Consumer

---

- Creating a `KafkaConsumer` is similar to creating a `KafkaProducer`
- Like the producer, you must specify `bootstrap.servers`,  
`key.deserializer`, and `value.deserializer` in a  
`Properties` object
- You must also specify a `group.id` which specifies the consumer group  
for which the `KafkaConsumer` instance belongs to

# Setup of KafkaConsumer Example

---

```
private Properties kafkaProps = new Properties();
kafkaProps.put("bootstrap.servers", "mybroker1:9092,mybroker2:9092");

kafkaProps.put("key.serializer",
"org.apache.kafka.common.serialization.StringSerializer");
kafkaProps.put("value.serializer",
"org.apache.kafka.common.serialization.StringSerializer");
kafkaProps.put("group.id", "StateCounter");

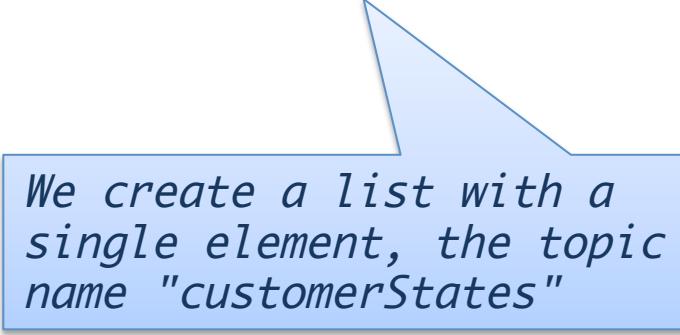
KafkaConsumer<String, String> consumer = new KafkaConsumer<String,
String>(kafkaProps);
```

# Subscribing to Topics

---

- Once a consumer is created, you can subscribe to one or more topics

```
consumer.subscribe(Collections.singletonList("customerStates"));
```



*We create a list with a single element, the topic name "customerStates"*

# Subscribing with Regular Expressions

---

- You can also subscribe to topics using regular expressions
- The expression can match multiple topic names
- If a new topic is created with a name that matches, a rebalance will happen and consumers will start consuming from the new topic
- Useful for applications that need to consume from multiple topics
- Example: subscribe to all test topics
  - `consumer.subscribe("test.*");`

# Consumer Poll Loop

---

- Once a consumer subscribes to topics, a loop polls the server for more data

```
KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props);
consumer.subscribe(Arrays.asList("foo", "bar"));
while (true) {
    ConsumerRecords<String, String> records = consumer.poll(100);

    for (ConsumerRecord<String, String> record : records)
        System.out.printf("offset = %d, key = %s, value = %s",
                          record.offset(),
                          record.key(),
                          record.value());
} finally {
    consumer.close();
}
```

# Consumer Poll Loop

- Once a consumer subscribes to topics, a loop polls the server for more data

```
KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props);
consumer.subscribe(Arrays.asList("foo", "bar"));
while (true) {
    ConsumerRecords<String, String> records = consumer.poll(100);

    for (ConsumerRecord<String, String> record : records)
        System.out.printf("offset = %d, key = %s, value = %s", record.offset(),
record.key(), record.value());
} finally {
    consumer.close();
}
```

*Consumers must keep polling Kafka or they will be considered dead and the partitions they are consuming will be handed to another consumer in the group*

# Poll Method

---

- `poll()` returns a list of records that contain:
  - Topic and partition the record came from
  - Offset of the record within the partition
  - Key and value of the record
- Takes a timeout parameter that specifies how long it will take to return, with or without data
  - *How fast do you want to return control to the thread that does the polling?*

# Poll Internals

---

- Behind the scenes of `poll()`:
  - First time `poll()` is called with a new consumer: finds the GroupCoordinator, joins the consumer group, and receives a partition assignment
  - If a rebalance is triggered, it is also handled inside the poll loop
  - Heartbeats that keep consumers alive are sent
- Processing between iterations must be fast and efficient

# Multithreading Considerations

---

- **One consumer per thread**
- To run multiple consumers in the same group in one application, each must run in its own thread
- You can wrap the consumer logic in its own object and then use Java's `ExecutorService` to start multiple threads each with its own consumer

# Commits

---

- Unlike other JMS queues, Kafka does not track acknowledgements from consumers
- When `poll()` is called, it returns records that consumers in the group have not yet read
- The records that have been read by a consumer of the group are tracked by their position (offset) in each partition
- When the current position in the partition is updated, it is known as a *commit*

# Consumers Commit Offsets

---

- Consumers send a message to Kafka to a reserved topic with the committed offset for each partition
- If a consumer crashes or a new consumer joins the consumer group, a rebalance is triggered
- After a rebalance, a consumer may be assigned a new set of partitions and must read the latest committed offset of each partition and continue from there

# Automatic Commit

---

- If you configure `enable.auto.commit=true`, the consumer will commit the largest offset your client received from `poll()` every 5 seconds by default
- Whenever there is a poll, the consumer checks if its time to commit and if so, will commit the offsets it returned in the last poll
- Convenient but don't give developers enough control to avoid duplicate messages

# Commit Current Offset

---

- Developers usually want to exercise control over the time offsets are committed to eliminate possibility of missing messages **and** reduce the number of duplicate messages during rebalancing
- Setting `auto.commit.offset=false` means that offsets will only be committed explicitly
- Consumer has a `commitSync()` API to commit the latest offset returned by `poll()`

# commitSync Example

```
while (true) {  
    ConsumerRecords<String, String> records = consumer.poll(100);  
  
    for (ConsumerRecord<String, String> record : records)  
        System.out.printf("offset = %d, key = %s, value = %s", record.offset(),  
    record.key(), record.value());  
  
    try {  
        consumer.commitSync();  
    } catch (CommitFailedException e) {  
        log.error("commit failed", e);  
    }  
  
} finally {  
    consumer.close();  
}
```

*Once we are done  
processing all records in  
the current batch,  
commitSync is called  
before polling for more*

# Asynchronous Commit

- The application is blocked until the broker responds to the commit request with `commitSync()` limiting throughput
- An alternative is to use `commitAsync()` which commits the last offsets and continues

```
while (true) {  
    ConsumerRecords<String, String> records = consumer.poll(100);  
  
    for (ConsumerRecord<String, String> record : records) {  
        System.out.printf("offset = %d, key = %s, value = %s%n",  
            record.offset(), record.key(), record.value());  
  
        consumer.commitAsync();  
    }  
}
```

*You can also pass a callback which is invoked by the consumer when the commit finishes (either successfully or not)*

# Summary

---

- Producers
  - API
  - Sending messages
  - Serialization
- Consumers
  - API
  - Subscribing
  - Consumer groups
  - Rebalance

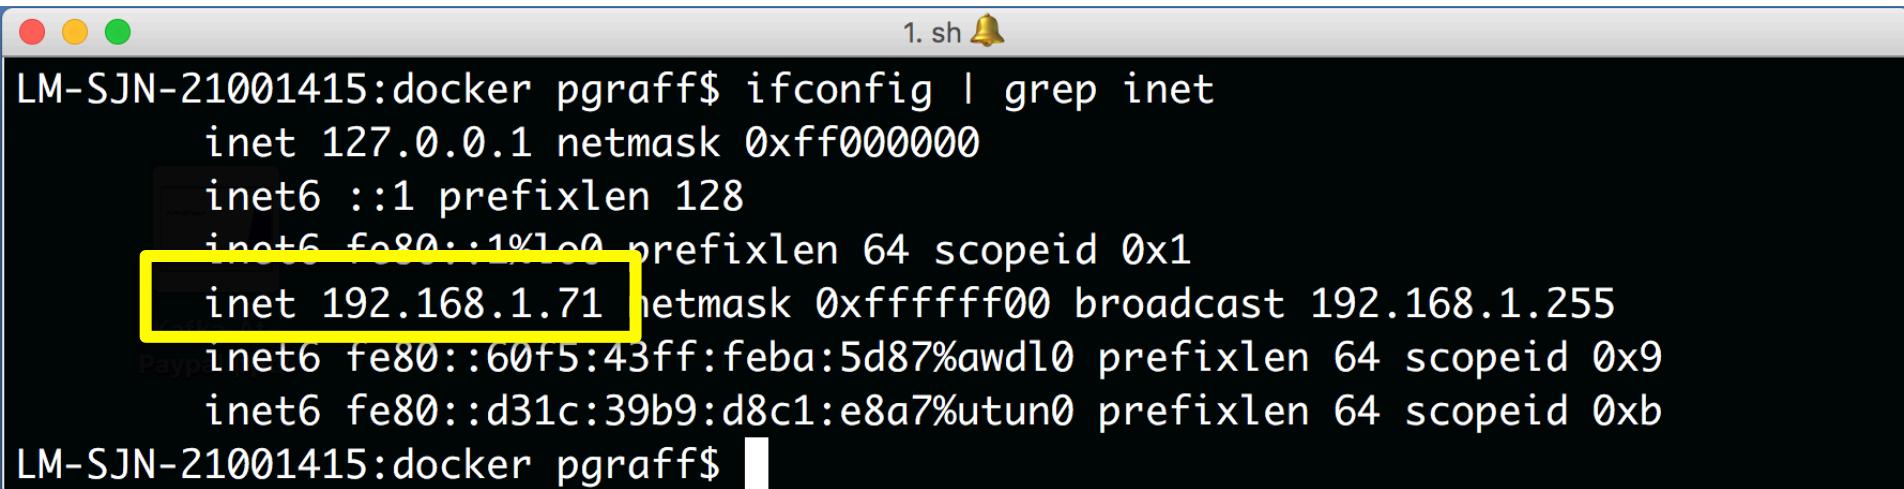
# Lab

---

- In this lab we'll use the Java API to consume and produce messages
- Two projects
  - Producer
  - Consumer
- Required tools:
  - Java
  - Maven
  - Docker (as before)
- Maven can be run stand alone or using a docker-container

# Step 1: Configure Your docker-compose.yml file

- We need to configure the docker-compose.yml with the IP of your machine
- To find the IP (**ifconfig** on the Mac, **ipconfig** on a Windows machine)



```
1. sh 🎙
LM-SJN-21001415:docker pgraff$ ifconfig | grep inet
    inet 127.0.0.1 netmask 0xff000000
    inet6 ::1 prefixlen 128
    inet6 fe80::1%lo0 prefixlen 64 scopeid 0x1
    inet 192.168.1.71 netmask 0xfffffff00 broadcast 192.168.1.255
    inet6 fe80::60f5:43ff:feba:5d87%awdl0 prefixlen 64 scopeid 0x9
    inet6 fe80::d31c:39b9:d8c1:e8a7%utun0 prefixlen 64 scopeid 0xb
LM-SJN-21001415:docker pgraff$ █
```

## Step 2: Edit your docker-compose.yml file

- Using your favorite editor, edit the docker-compose.yml file in the directory:  
**kafka-labs/labs/02-Publish-And-Subscribe/docker**
- Look for the [insert] statement and replace it with your IP... E.g.:

```
- 9092:9092  
- 7203:7203  
environment:  
  KAFKA_ADVERTISED_HOST_NAME: 192.168.1.70  
  KAFKA_ADVERTISED_PORT: 9092  
  KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181  
depends_on:
```

# Step 3: Run the docker-compose command

- In the docker directory (where you just modified the docker-compose file), run:  
**docker-compose up**

```
LM-SJN-21001415:docker pgraff$ docker-compose up
Creating network "docker_default" with the default driver
Creating docker_zookeeper_1 ...
Creating docker_zookeeper_1 ... done
Creating docker_kafka_1 ...
Creating docker_kafka_1 ... done
Attaching to docker_zookeeper_1, docker_kafka_1
zookeeper_1  | ZooKeeper JMX enabled by default
zookeeper_1  | Using config: /conf/zoo.cfg
kafka_1      | waiting for kafka to be ready
zookeeper_1  | 2017-11-11 23:56:03,305 [myid:] - INFO  [main:QuorumPeerConfig@124] - Ready
```

# Step 4: Create the Topics

---

- In a second terminal, cd into the directory  
**kafka-lab/labs/02-Publish-And-Subscribe/docker**
- Create the two topics required for this exercise (user-event, global-events)
  - user-events
  - global-events

# Step 5: Build the Producer

- Use maven (either natively or using the docker image as described in the exercise description on GitHub)
- Build using mvn package from the directory  
**kafka-lab/labs/02-Publish-And-Subscribe/producer**

```
sh
LM-SJN-21001415:producer pgraff$ mvn package
[INFO] Scanning for projects...
[INFO]
[INFO] -----
[INFO] Building pubsub-producer 1.0-SNAPSHOT
[INFO] -----
[INFO]
[INFO] --- maven-resources-plugin:2.6:resources (default-resources) @ pubsub-producer ---
[WARNING] Using platform encoding (UTF-8 actually) to copy filtered resources, i.e. build is platform dependent!
[INFO] Copying 1 resource
[INFO]
[INFO] --- maven-compiler-plugin:3.5.1:compile (default-compile) @ pubsub-producer ---
[INFO] Nothing to compile - all classes are up to date
[INFO]
[INFO] --- maven-resources-plugin:2.6:testResources (default-testResources) @ pubsub-producer ---
[WARNING] Using platform encoding (UTF-8 actually) to copy filtered resources, i.e. build is platform dependent!
```

# Producer code

```
1 for (int i = 0; i < 100000; i++) {  
2     producer.send(new ProducerRecord<String, String>(  
3         "user-events", // topic  
4         "user_id_" + i, // key  
5         "some_value_" + System.nanoTime())); // value  
6  
7     if (i % 100 == 0) {  
8         String event = global_events[(int)(Math.random() * global_events.  
9  
10    producer.send(new ProducerRecord<String, String>(  
11        "global-events", // topic  
12        event)); // value  
13  
14    producer.flush();  
15    System.out.println("Sent message number " + i);  
16 }
```

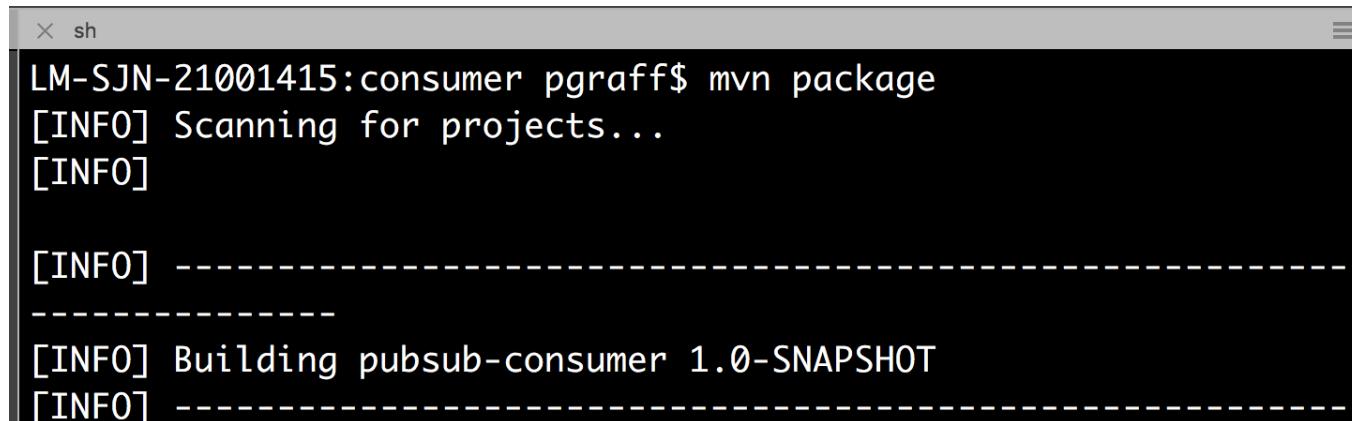
# Step 6: Run the Publisher

- The maven job creates a 'real executable' (target/producer)
- We'll run the executable

```
X sh
LM-SJN-21001415:producer pgraff$ target/producer
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
Sent message number 0
Sent message number 100
Sent message number 200
Sent message number 300
Sent message number 400
Sent message number 500
Sent message number 600
Sent message number 700
Sent message number 800
Sent message number 900
```

# Step 7: Build the Consumer

- Open another terminal (just so that we can run the producer again if we want to)
- Change into the directory  
**kafka-lab/labs/02-Publish-And-Subscribe/consumer**
- Run the maven command to build:

A screenshot of a terminal window titled "sh". The window contains the following text:

```
LM-SJN-21001415:consumer pgraff$ mvn package
[INFO] Scanning for projects...
[INFO]

[INFO] -----
[INFO] -----
[INFO] Building pubsub-consumer 1.0-SNAPSHOT
[INFO] -----
```

The text is white on a black background, with some lines wrapped to fit the window.

# The Consumer...

```
13 class Consumer {
14     public static void main(String[] args) throws IOException {
15         KafkaConsumer<String, String> consumer;
16         try (InputStream props = Resources.getResource("consumer.properties").openStream()) {
17             Properties properties = new Properties();
18             properties.load(props);
19             consumer = new KafkaConsumer<>(properties);
20         }
21         consumer.subscribe(Arrays.asList("user-events", "global-events"));
22
23     while (true) {
24         ConsumerRecords<String, String> records = consumer.poll(100);
25
26         for (ConsumerRecord<String, String> record : records) {
27             switch (record.topic()) {
28                 case "user-events":
29                     System.out.println("Received user-events message - key: " + record.key() + " value
30                     break;
31                 case "global-events":
32                     System.out.println("Received global-events message - value: " + record.value());
33                     break;
34                 default:
35                     throw new IllegalStateException("Shouldn't be possible to get message on topic " +
36                 }
37             }
38 }
```

# Step 8: Run the consumer

- Run the consumer (it will be able to consume the already produced messages as it is 'playing from the beginning')
- The consumer is in **target/consumer**

```
java
Received user-events message - key: user_id_99982 value: some_value_125646276189312
Received user-events message - key: user_id_99983 value: some_value_125646276190697
Received user-events message - key: user_id_99984 value: some_value_125646276192110
Received user-events message - key: user_id_99985 value: some_value_125646276193473
Received user-events message - key: user_id_99986 value: some_value_125646276194923
Received user-events message - key: user_id_99987 value: some_value_125646276196349
Received user-events message - key: user_id_99988 value: some_value_125646276197731
Received user-events message - key: user_id_99989 value: some_value_125646276199341
Received user-events message - key: user_id_99990 value: some_value_125646276200693
Received user-events message - key: user_id_99991 value: some_value_125646276202068
Received user-events message - key: user_id_99992 value: some_value_125646276203454
Received user-events message - key: user_id_99993 value: some_value_125646276204790
Received user-events message - key: user_id_99994 value: some_value_125646276206252
Received user-events message - key: user_id_99995 value: some_value_125646276207581
Received user-events message - key: user_id_99996 value: some_value_125646276209042
Received user-events message - key: user_id_99997 value: some_value_125646276210568
Received user-events message - key: user_id_99998 value: some_value_125646276211931
Received user-events message - key: user_id_99999 value: some_value_125646276213364
```

× docker-compose

```
consumer_offsets,45] (kafka.coordinator.GroupMetadataManager)
kafka_1    | [2017-11-12 02:07:47,257] INFO [Group Metadata Manager on Broker 1001]: Finished loading offsets from [__consumer_
ffsets,45] in 1 milliseconds. (kafka.coordinator.GroupMetadataManager)
kafka_1    | [2017-11-12 02:07:47,158] INFO [Group Metadata Manager on Broker 1001]: Finished loading offsets and group metadata from [__consumer_
consumer_offsets,48] (kafka.coordinator.GroupMetadataManager)
kafka_1    | [2017-11-12 02:07:47,259] INFO [Group Metadata Manager on Broker 1001]: Finished loading offsets from [__consumer_
ffsets,48] in 1 milliseconds. (kafka.coordinator.GroupMetadataManager)
kafka_1    | [2017-11-12 02:07:47,305] INFO [GroupCoordinator 1001]: Preparing to restabilize group test with old generation 0
kafka.coordinator.GroupCoordinator
kafka_1    | [2017-11-12 02:07:47,314] INFO [GroupCoordinator 1001]: Received assignment for group test for generation 0 from kafka.coordinator.Group
Coordinator)
kafka_1    | [2017-11-12 02:07:47,321] INFO [GroupCoordinator 1001]: Assignment received from leader for group test for generat
on 1 (kafka.coordinator.GroupCoordinator)
```

# Producer

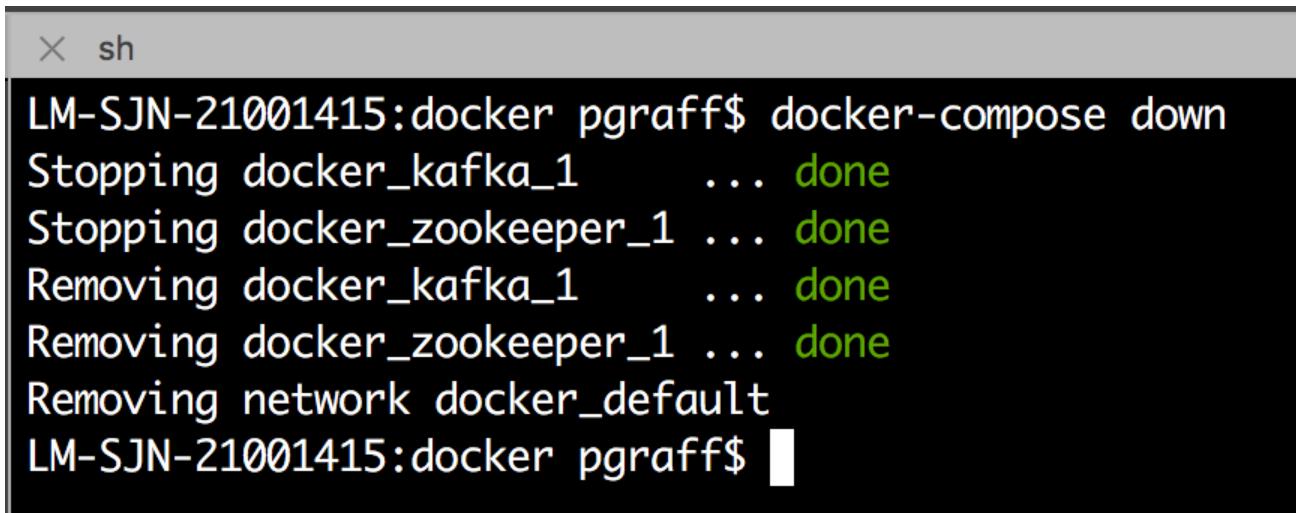
```
Sent message number 98200
Sent message number 98300
Sent message number 98400
Sent message number 98500
Sent message number 98600
Sent message number 98700
Sent message number 98800
Sent message number 98900
Sent message number 99000
Sent message number 99100
Sent message number 99200
Sent message number 99300
Sent message number 99400
Sent message number 99500
Sent message number 99600
Sent message number 99700
Sent message number 99800
Sent message number 99900
```

# Consumer

```
7833210
Received user-events message - key: user_id_99989 value: some_value_1261761
7834601
Received user-events message - key: user_id_99990 value: some_value_1261761
7835831
Received user-events message - key: user_id_99991 value: some_value_1261761
7837097
Received user-events message - key: user_id_99992 value: some_value_1261761
7838306
Received user-events message - key: user_id_99993 value: some_value_1261761
7839570
Received user-events message - key: user_id_99994 value: some_value_1261761
7843160
Received user-events message - key: user_id_99995 value: some_value_1261761
7844334
Received user-events message - key: user_id_99996 value: some_value_1261761
7845470
Received user-events message - key: user_id_99997 value: some_value_1261761
7846637
Received user-events message - key: user_id_99998 value: some_value_1261761
7847738
Received user-events message - key: user_id_99999 value: some_value_1261761
7851492
Received global-events message - value: plan_removed_126176143550296
Received global-events message - value: plan_added_126176145745964
```

# When You're Done...

- Remember to run the docker-compose down



```
sh
LM-SJN-21001415:docker pgraff$ docker-compose down
Stopping docker_kafka_1    ... done
Stopping docker_zookeeper_1 ... done
Removing docker_kafka_1    ... done
Removing docker_zookeeper_1 ... done
Removing network docker_default
LM-SJN-21001415:docker pgraff$
```