

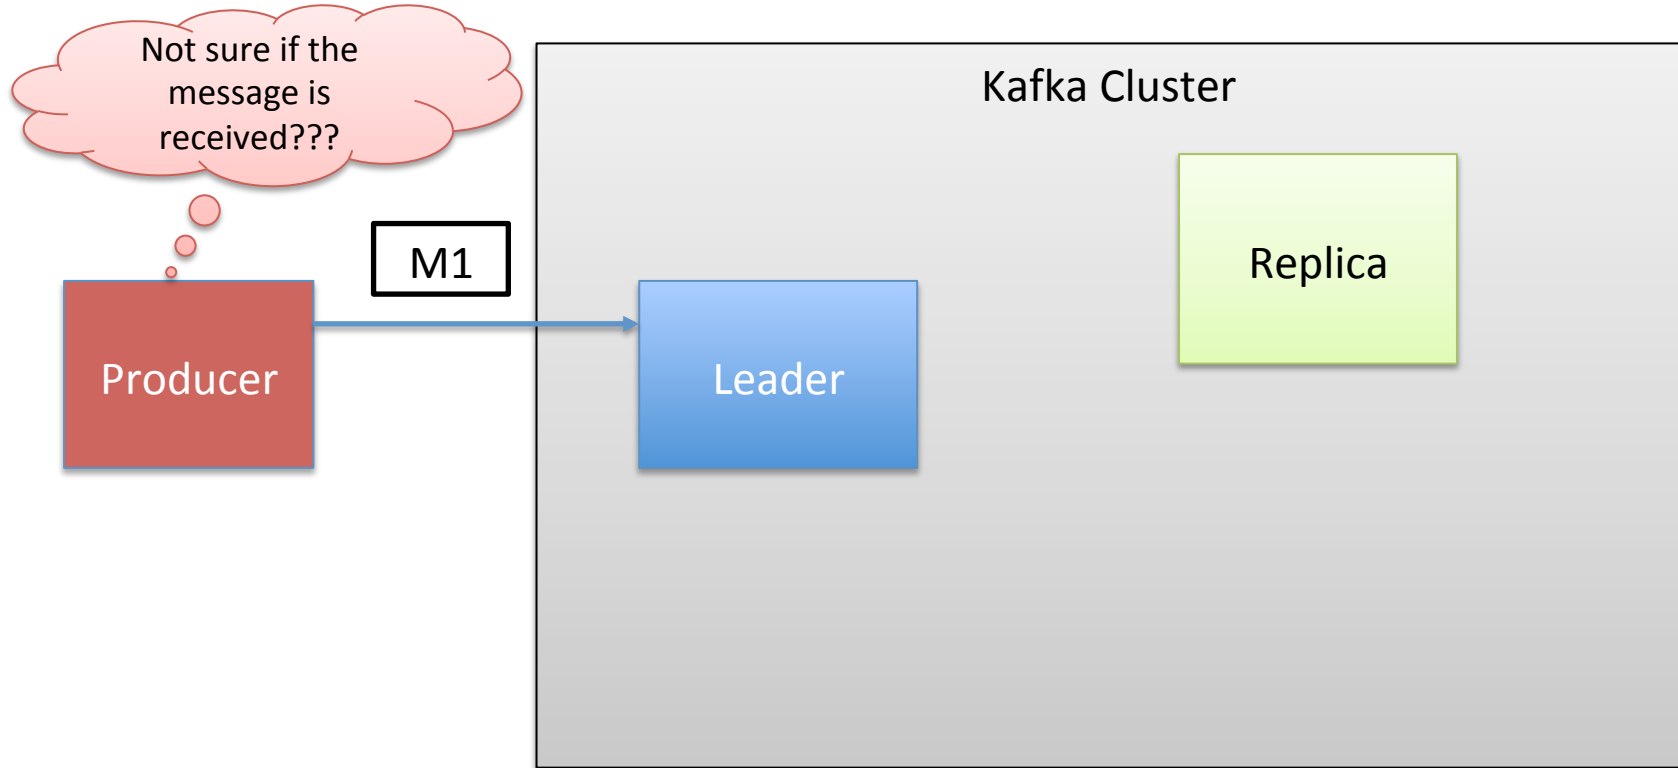
# Kafka Delivery Guarantees

# Introduction

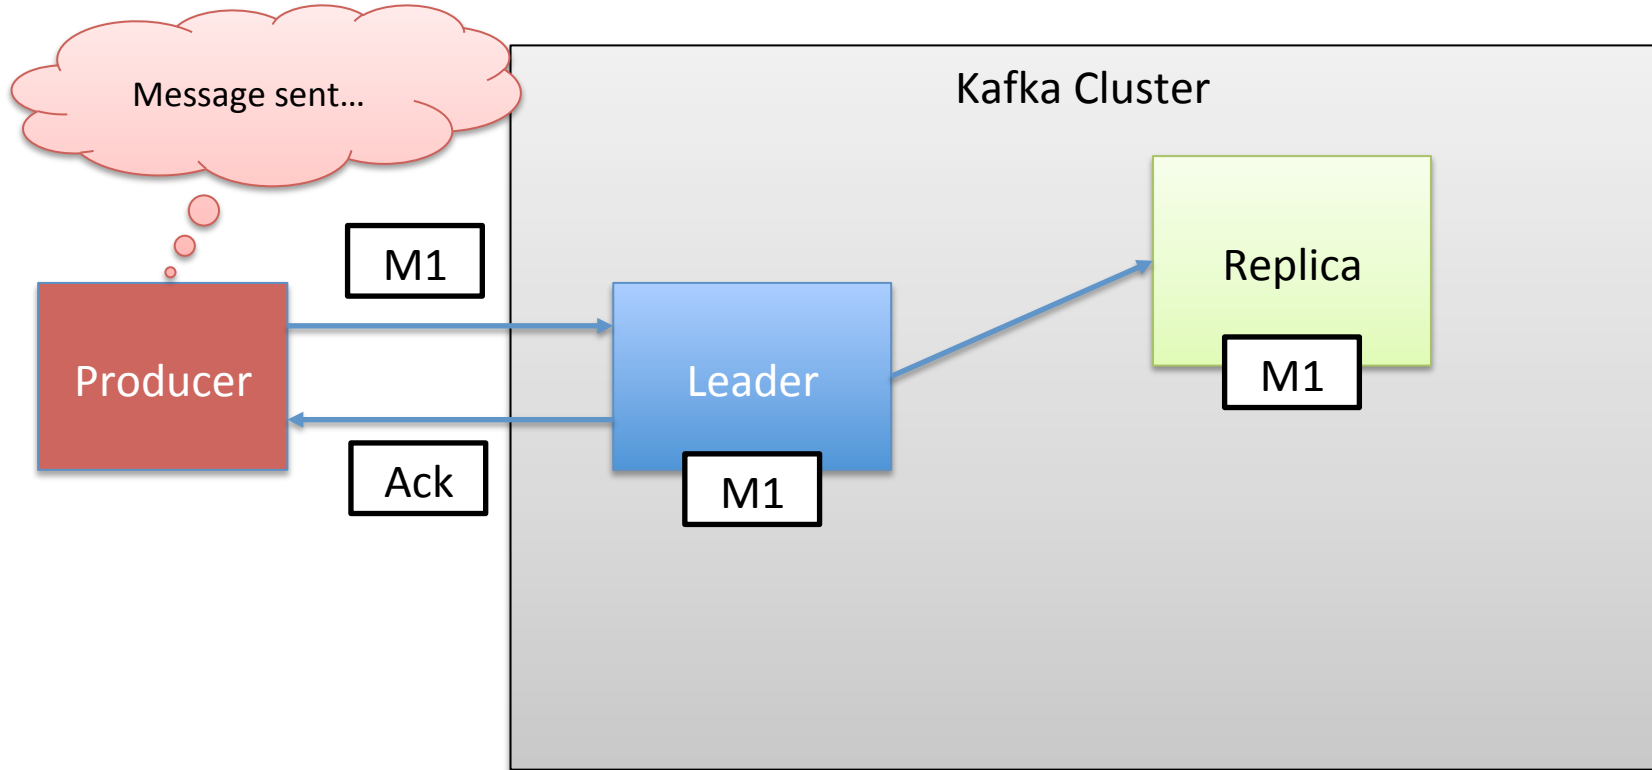
---

- Kafka's delivery model leaves potential gaps in typical enterprise deployment scenarios
- Naïve programming may lead to
  - Messages being processed multiple times
  - Messages not being processed
- We will look at a few different scenarios where such failures may occur
- We'll also look at some technique that can be used to minimize or avoid the failures
  - Idempotent services
  - Kafka's new exactly once processing guarantee

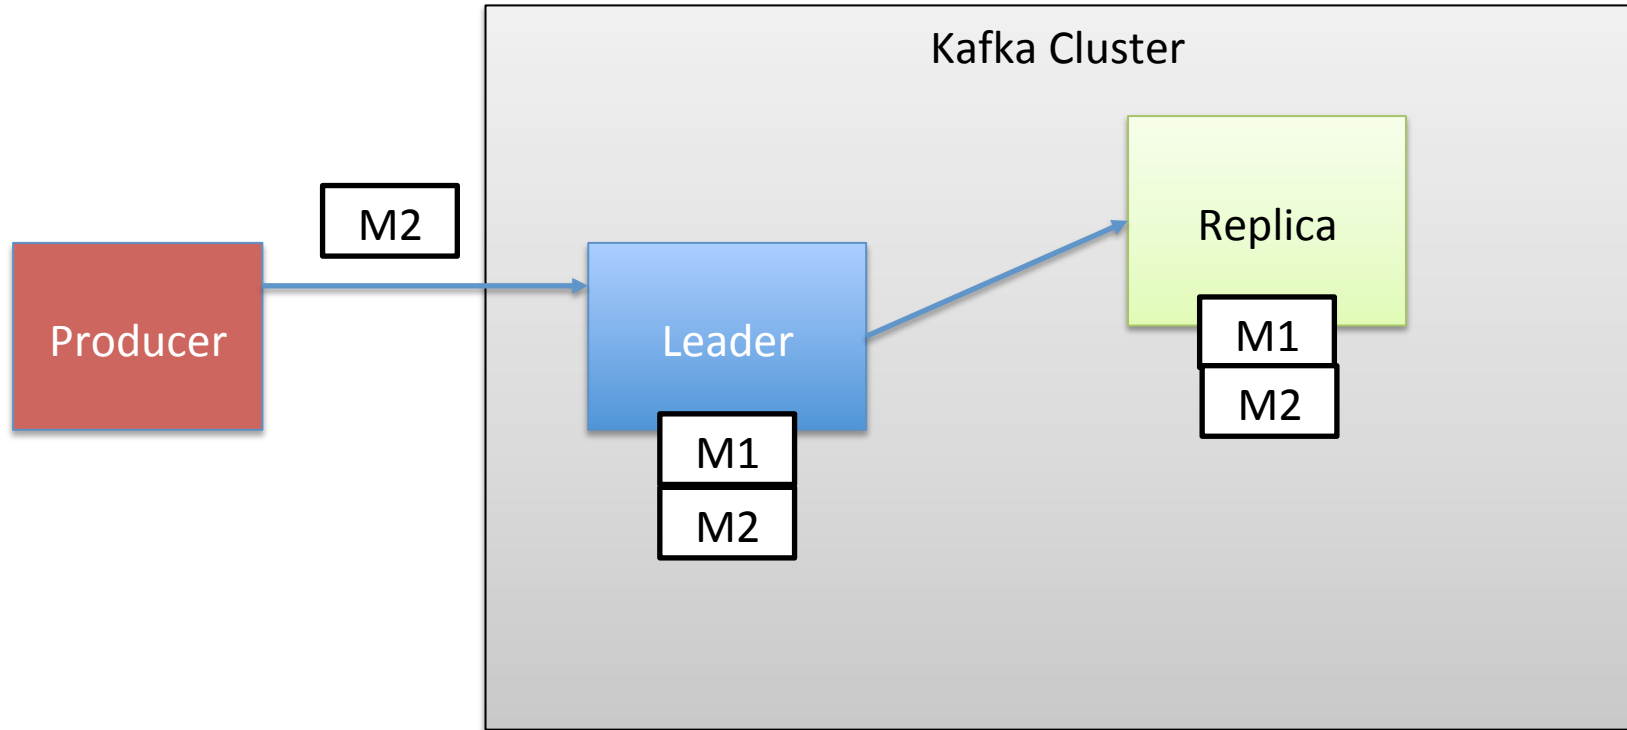
# Message Duplication: 1. Send the message



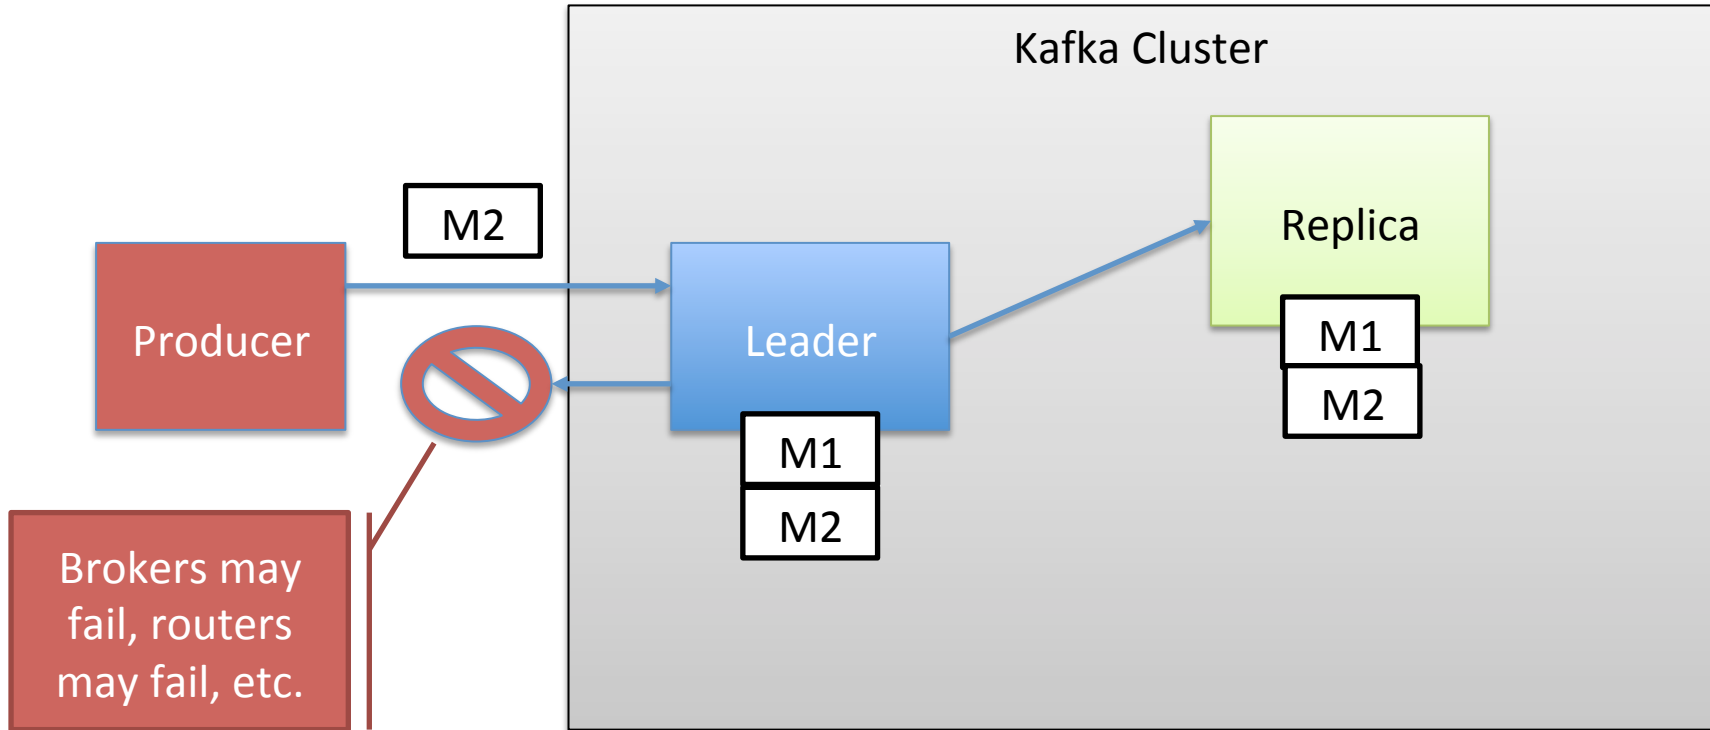
## Message Duplication: 2. Message ack



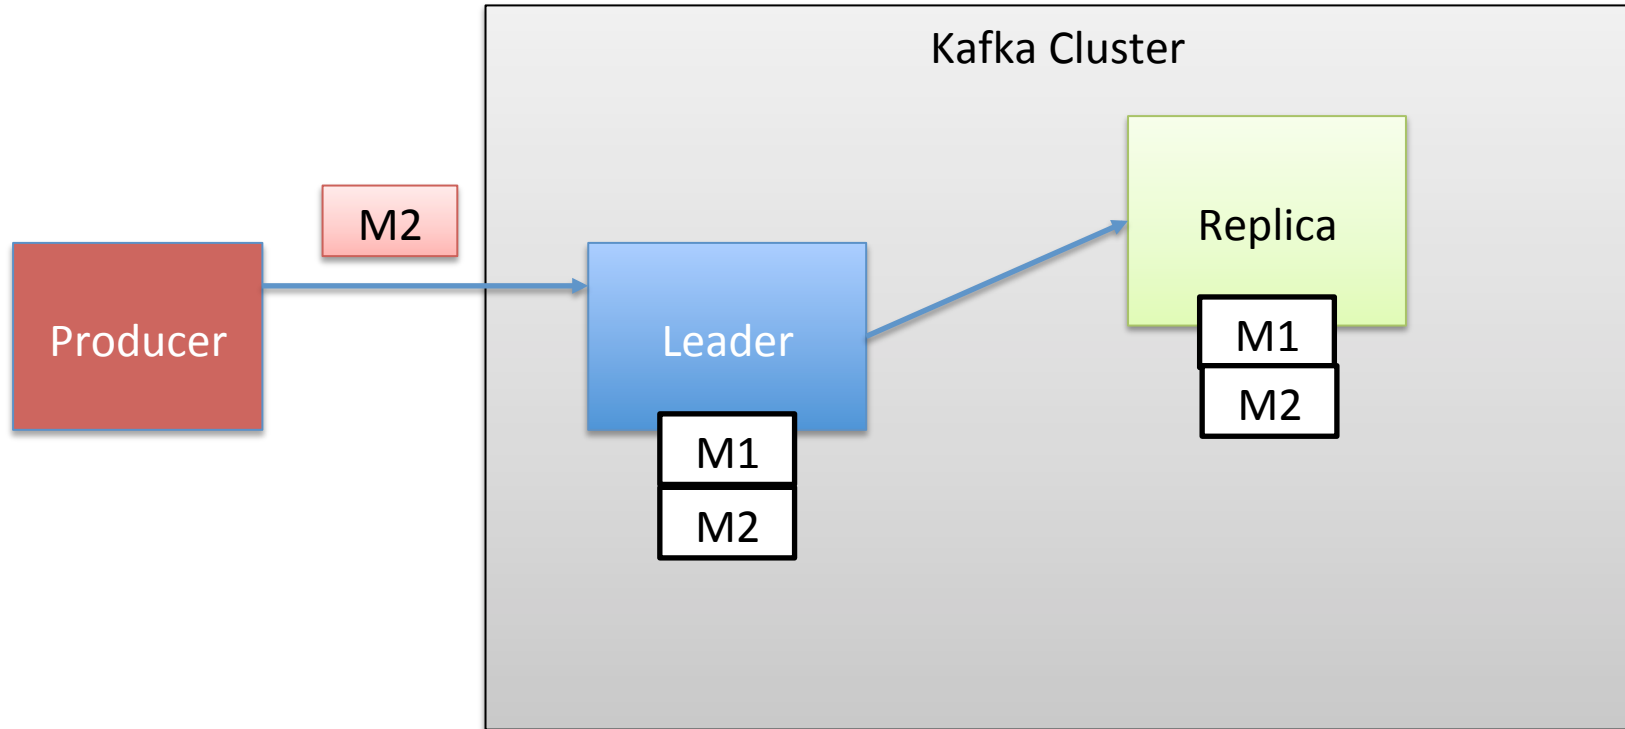
# Message Duplication: 3. M2 being sent



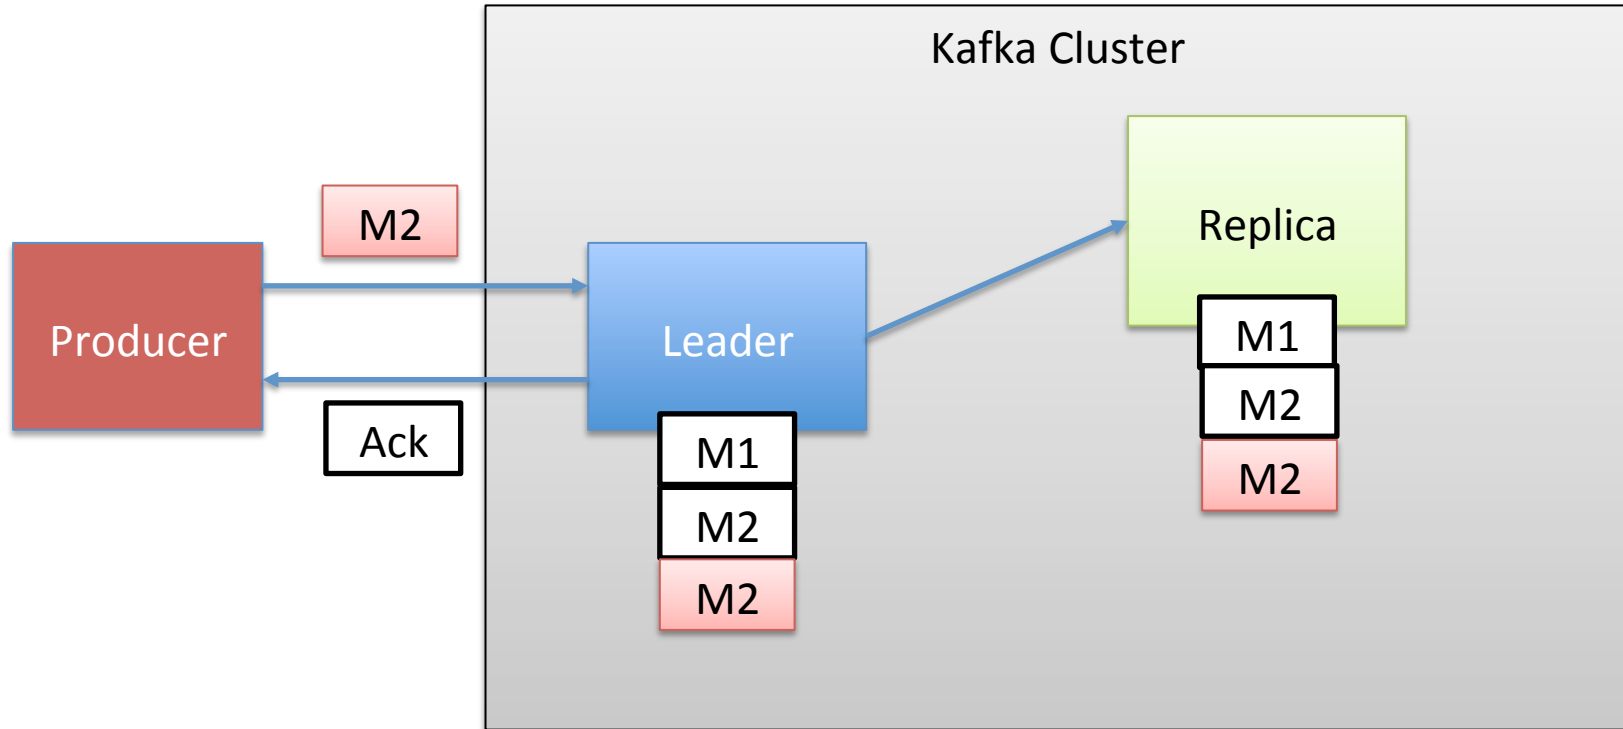
# Message Duplication: 4. Message never ack-ed



# Message Duplication: 5. M2 is resent



# Message Duplication: 5. M2 is resent and stored again



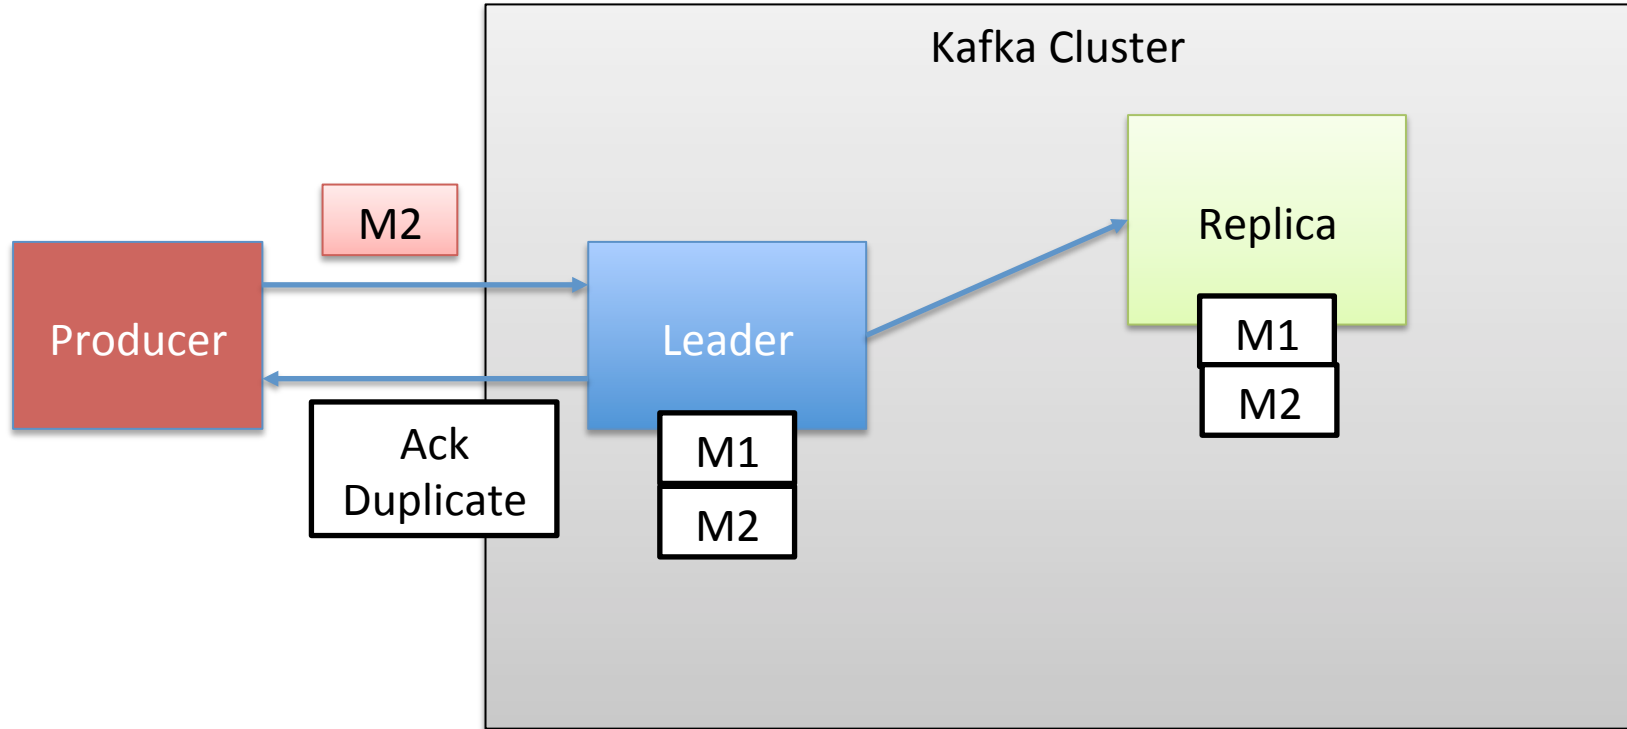


# Kafka Exactly Once Producer Solution

---

- Idempotent producer (per partition)
  - Exactly once
  - In order
- Transactions
  - Atomic writes across partitions

# Idempotent Producer



# Transaction API

---

- Atomic writes across multiple partitions
- Either all records are visible, or none

```
producer.initTransaction();
try {
    producer.beginTransaction();
    producer.send(record1);
    producer.send(record2);
    ...
    producer.commitTransaction();
}
catch (KafkaException e) {
    producer.abortTransaction();
}
```

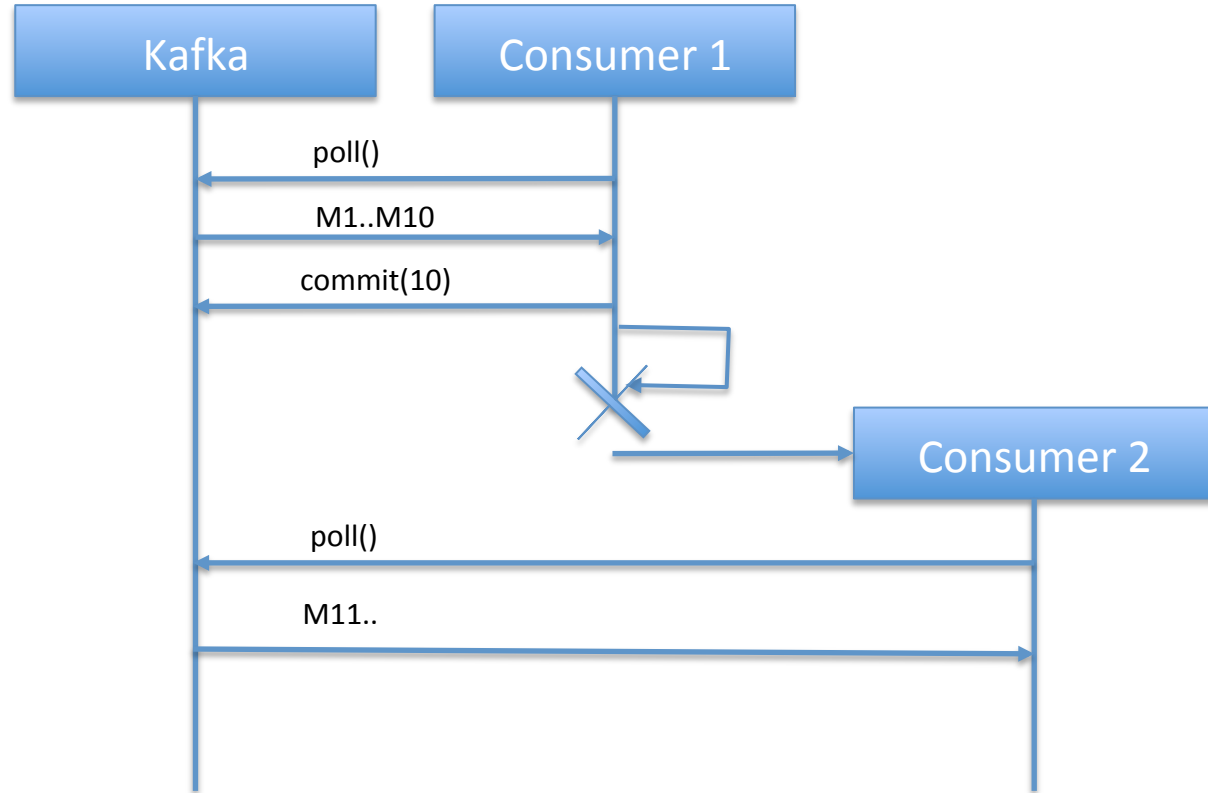
# Consumer Problems

---

- Messages missed
  - Offset is committed prior to processing the messages
  - No or only partial processing of incoming messages
  - Consumer failure
- Messages double processed
  - Processing of messages prior to commit
  - Consumer failure before commit
- There is now real failsafe way to ensure that messages are processed exactly once
  - Requires distributed transaction

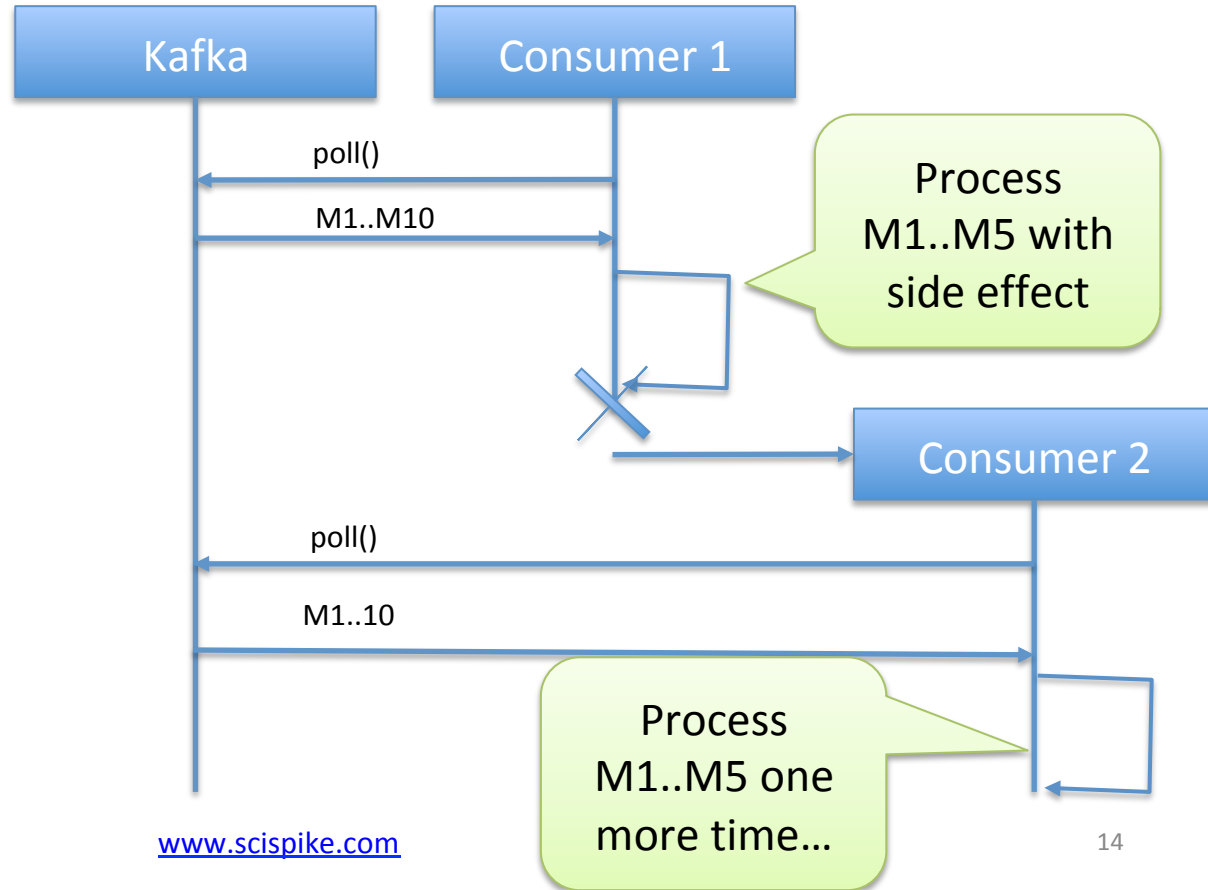
# Consumer Message Loss

1. Messages received (M1...M10)
2. Offset committed (M10)
3. Processing M1..M5
4. Consumer crashes
5. Consumer resumes
6. Message M11.. is delivered



# Consumer Double Processing

1. Messages received (M1...M10)
2. Processing M1..M5
3. Consumer crashes
4. Consumer resumes
5. Message M11.. is delivered



# Kafka Exactly Once Stream Processing

---

- Kafka provides an exactly once guarantee for stream processing
- Important to understand the guarantee
  - The effect on the Kafka stream state is AS IF each message was processed exactly once
    - The message may actually be processed multiple times
    - However, the offset is moved with the state changes
- This means:
  - If the processing have side effect OUTSIDE KAFKA this side effect may have to be idempotent!

# Configuration of Exactly Once

---

## Producer Configuration

`enable.idempotence=true`

### **Also recommended**

`max.inflight.requests.per.connection=1`

`acks="all"`

`retries= MAX_INT`

## Consumer Configuration

`isolation.level = read_committed`

or

`isolation.level = read_uncommitted`

`processing.mode = "exactly_once"`



# Summary

---

- Kafka supports an exactly once message delivery guarantee
- Solved with
  - Idempotent producer
  - Transactions
  - Stream processing
- It's important to know that the exactly once guarantee on the consumer is only guaranteed within the context of Kafka!